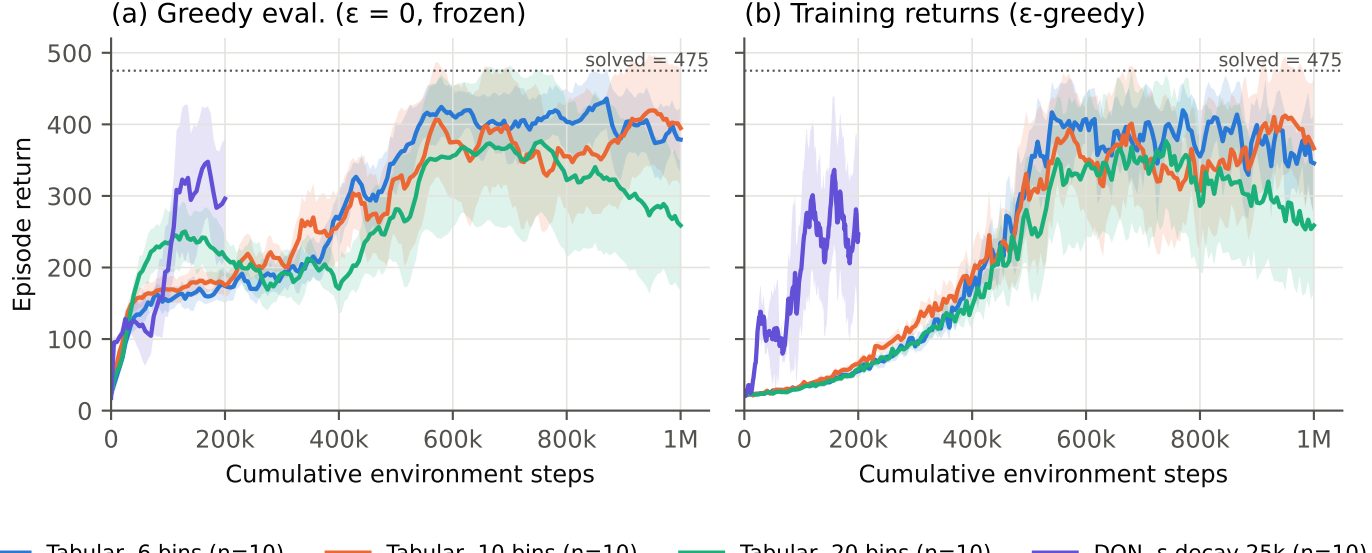
(a) Greedy eval. (ε = 0, frozen)
(b) Training returns (ε-greedy)
Episode return
Cumulative environment steps
solved = 475
0
100
200
300
400
500
200k
400k
600k
800k
1M
Tabular, 6 bins (n=10)
Tabular, 10 bins (n=10)
Tabular, 20 bins (n=10)
DQN, ε decay 25k (n=10)