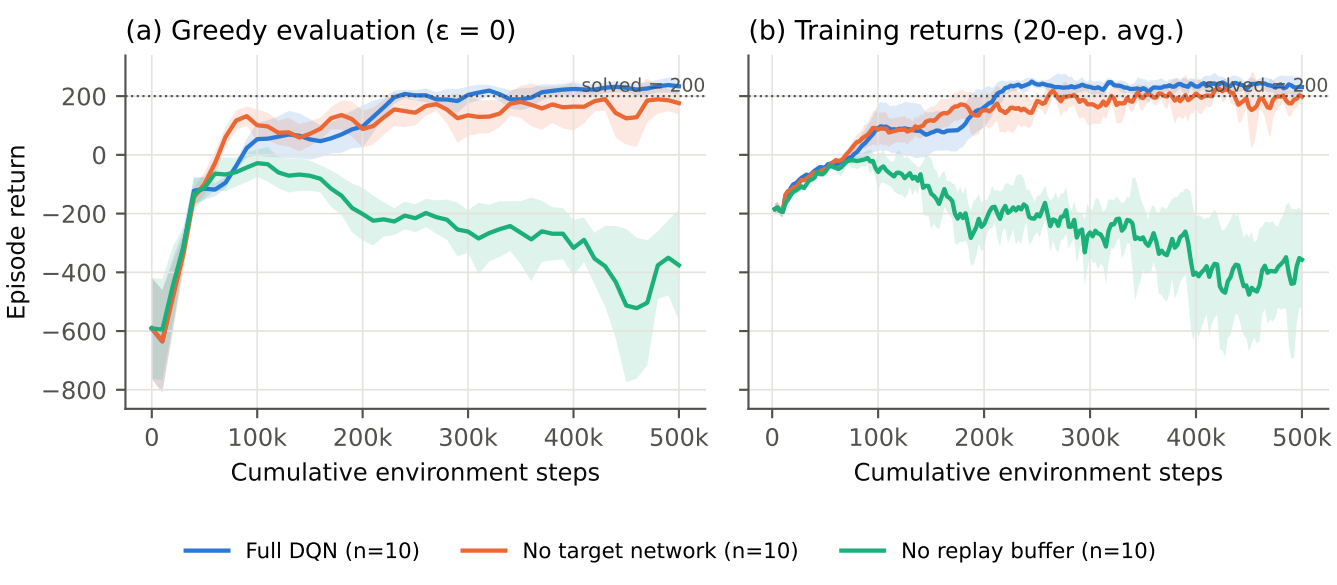
(a) Greedy evaluation (ε = 0)
(b) Training returns (20-ep. avg.)
Episode return
200
0
−200
−400
−600
−800
solved = 200
0
100k
200k
300k
400k
500k
Cumulative environment steps
Full DQN (n=10)
No target network (n=10)
No replay buffer (n=10)