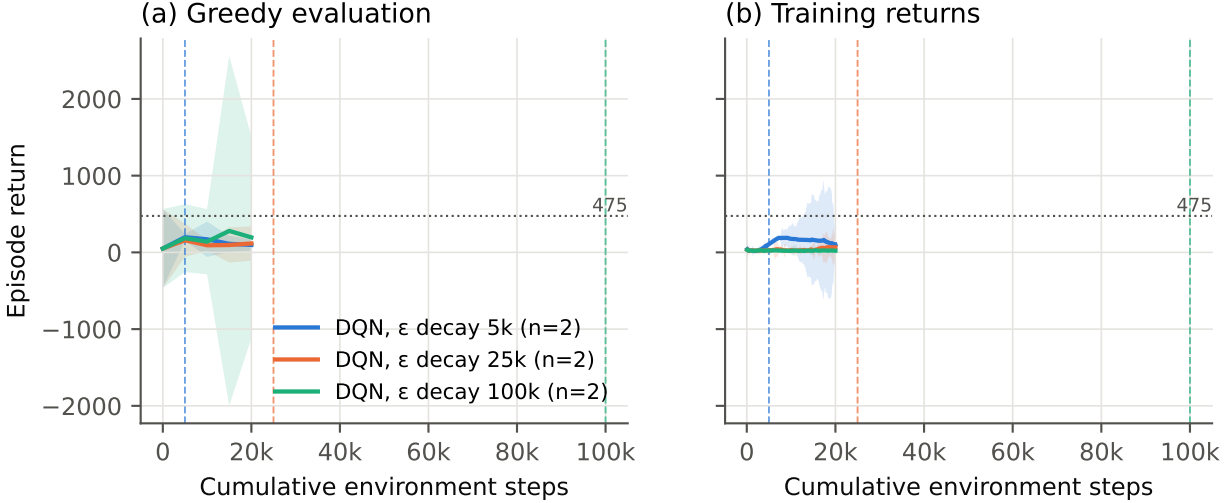
(a) Greedy evaluation
(b) Training returns
Episode return
2000
1000
0
−1000
−2000
475
DQN, ε decay 5k (n=2)
DQN, ε decay 25k (n=2)
DQN, ε decay 100k (n=2)
0
20k
40k
60k
80k
100k
Cumulative environment steps