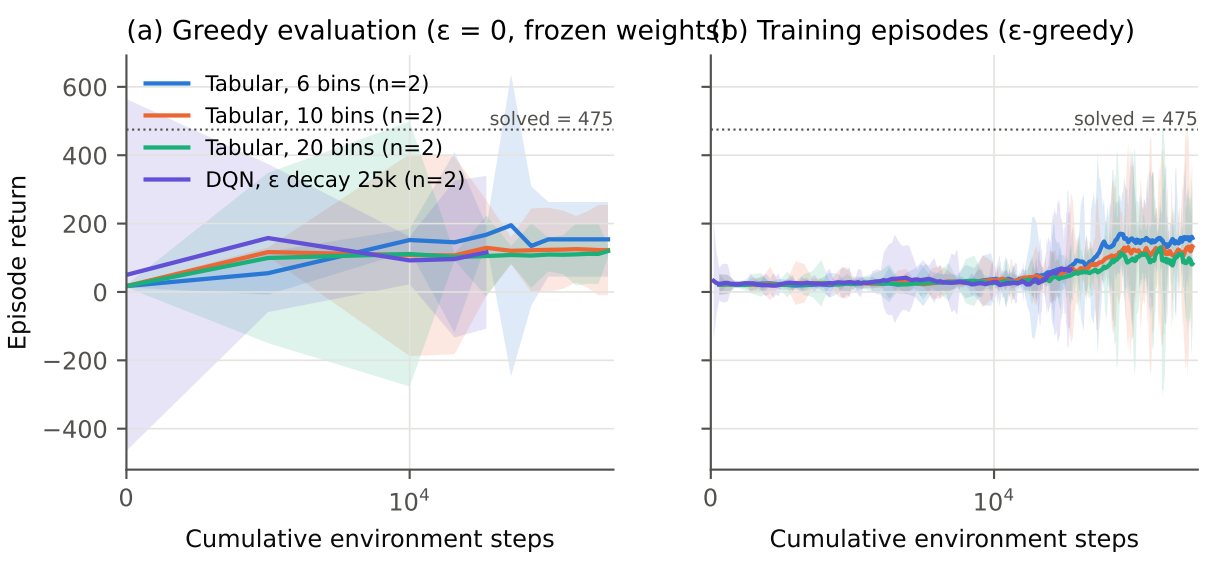
(a) Greedy evaluation (ε = 0, frozen weights)
(b) Training episodes (ε-greedy)
Tabular, 6 bins (n=2)
Tabular, 10 bins (n=2)
Tabular, 20 bins (n=2)
DQN, ε decay 25k (n=2)
solved = 475
solved = 475
Episode return
600
400
200
0
−200
−400
0
10^4
0
10^4
Cumulative environment steps
Cumulative environment steps